

TiDB Workshop

Friday 28 July, Berlin, WeAreDevelopers

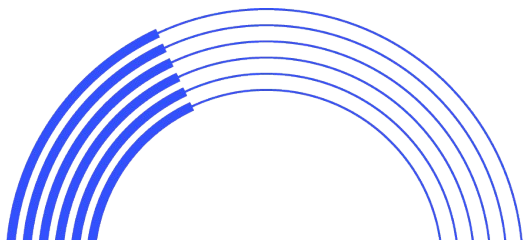


Table of Content

1. Introduction
2. Environment setup
3. Lab exercises
4. Conclusions

Introduction

- Mattias Jonsson
 - Working for PingCAP as Senior Database Engineer, developing TiDB
 - Previous:
 - Senior Developer / Engineering manager at Booking.com
 - Senior Software Engineer at MySQL/Sun/Oracle
- Daniël van Eeden
 - Working for PingCAP as Technical Support Engineer
 - Previous
 - Senior Database Engineer at Booking.com
 - Long time member in the MySQL Community

About PingCAP

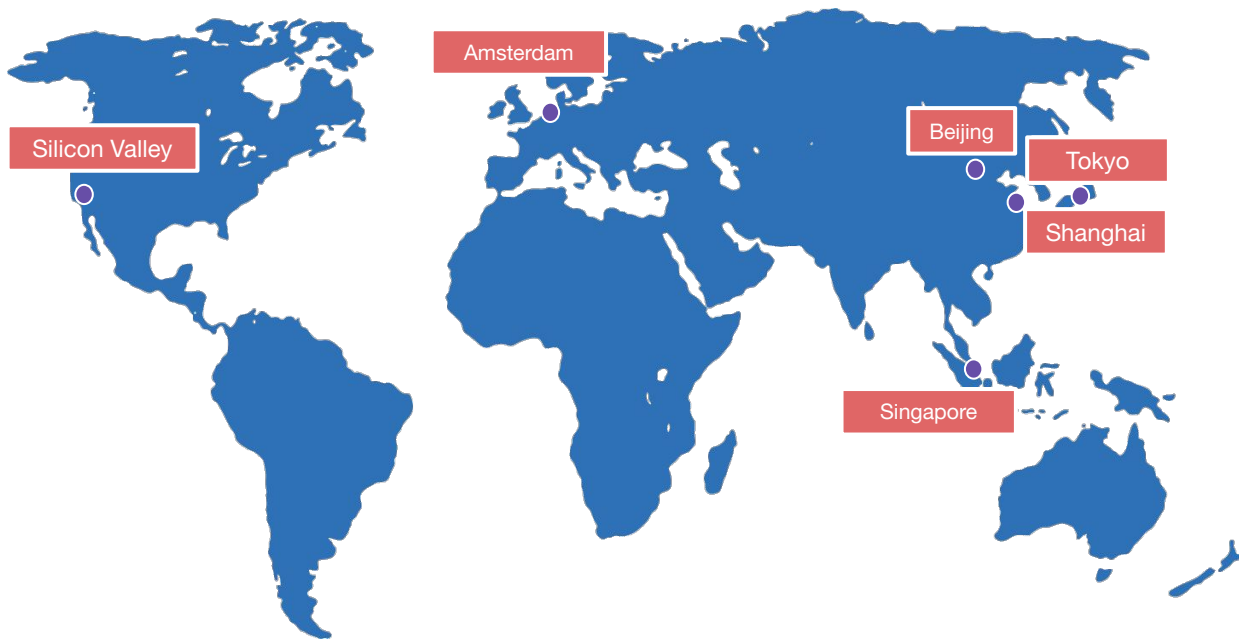
PingCAP, the company behind **TiDB**, provides enterprise-level and cloud-based services and technology for TiDB, so companies can count on TiDB as its core database solution to simplify the database infrastructure and create business value faster than ever.

Mission: Empower engineers to innovate with Speed, Agility, and Scale.



About PingCAP

- Founded in 2015
- The company behind TiDB
- Global
- Open source culture
- Strong investors
- 600+ employees



About PingCAP

Trusted and Verified by Vertical Leaders Globally

2000+ production adopters

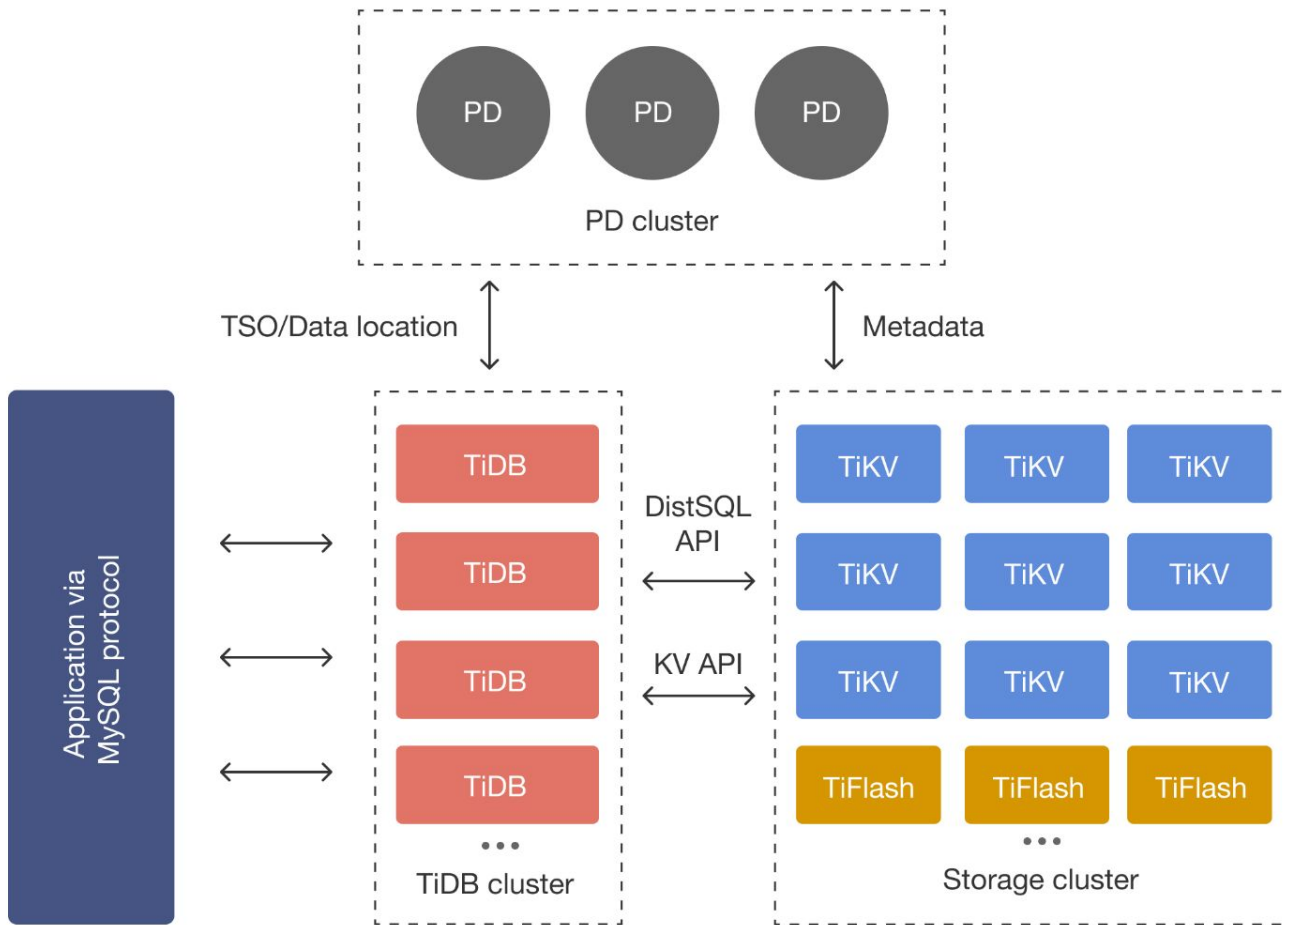


And

Airbnb, Pinterest, Niantic

2 out of top 3 Cryptocurrency Exchanges

One of Top 3 biggest banks all over the world



Environment setup

Database:

- TiDB Locally with TiUP via <https://tiup.io>
- MySQL Locally
- Optional: TiDB Cloud
 - via <https://tidbcloud.com> (Free serverless tier)
 - select eu-central-1 as location for your cluster

Application:

- demoblog
- <https://github.com/dveeden/demoblog>

Database client:

- MySQL Client
- TiDB Cloud Console (Chat2Query)

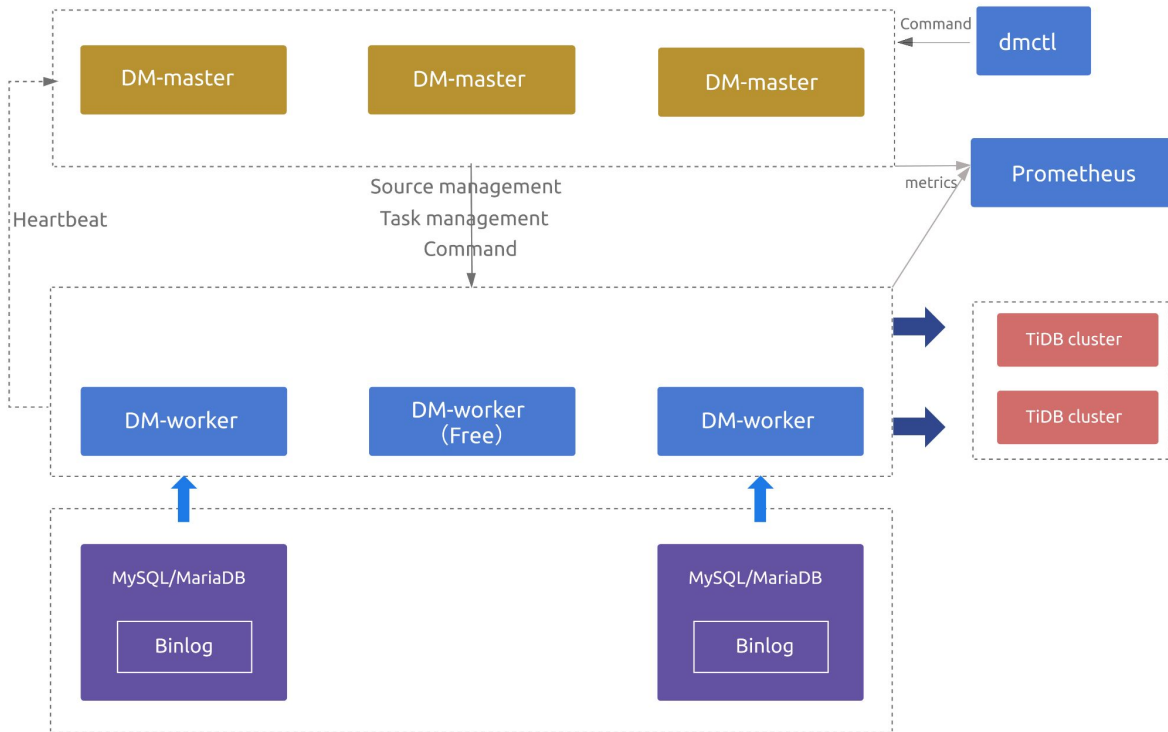
Environment setup

- AWS
- This presentation (for easier li
- EC2 hosts, pick one from [this spreadsheet](#).
- shared keypair
 - For simplicity we use the same private key for ssh
 - Log in with your ssh client, like:
`ssh -i tidb-workshop.pem ubuntu@3.78.221.19`
 - Download the demoblog repository:
`git clone https://github.com/dveeden/demoblog.git`
 - Setup the environment (will also start MySQL Server)
`./demoblog/AWS/setup_workshop.sh`
- port forwarding (for later)
 - `ssh -L8080:127.0.0.1:8080 -L2379:127.0.0.1:2379 -L3000:127.0.0.1:3000 -i tidb-workshop.pem ubuntu@3.78.221.19`
 - <http://127.0.0.1:8080> Demoblog
 - <http://127.0.0.1:2379/dashboard> PD dashboard (root/)
 - <http://127.0.0.1:3000/> Grafana (admin/admin)

Migrate to TiDB from MySQL

Data Migration (DM)

- Replicate data from MySQL to TiDB.
- Uses Dumpling/Lightning to copy the initial copy.
- Note that this is also made to be high available.



Lab excercises: DM 1/2

1. Run demoblog with MySQL

```
~/demoblog/demoblog -db 'blog:blog@tcp(127.0.0.1)/blog?parseTime=true'
```

2. Make sure the application works (check <http://127.0.0.1:8080> (via portforward))

3. Run TiUP Playground (in a new ssh session)

```
tiup playground v7.2.0
```

4. Run DM master (in a new ssh session)

```
./dm-master -master-addr 127.0.0.1:8261
```

5. Run DM worker (in a new ssh session)

```
./dm-worker -worker-addr 127.0.0.1:8262 -join 127.0.0.1:8261
```

6. Check DM status (in a new ssh session)

```
./dmctl -master-addr 127.0.0.1:8261 query-status
```

Lab exercises: DM 2/2

Repo:

- DM/mysql-source.yaml
- DM/dm-task.yaml

```
cd demoblog/DM
~/dmctl -master-addr 127.0.0.1:8261 operate-source create mysql-source.yaml
~/dmctl -master-addr 127.0.0.1:8261 start-task dm-task.yaml
~/dmctl -master-addr 127.0.0.1:8261 query-status mysql-to-tidb
```

Check:

```
sudo mysql blog
mysql -u root -h 127.0.0.1 -P 4000 blog
```

Query: SELECT * FROM posts;

Stop application and run again with TiDB (go back to the demoblog session, to simulate a switch to TiDB)

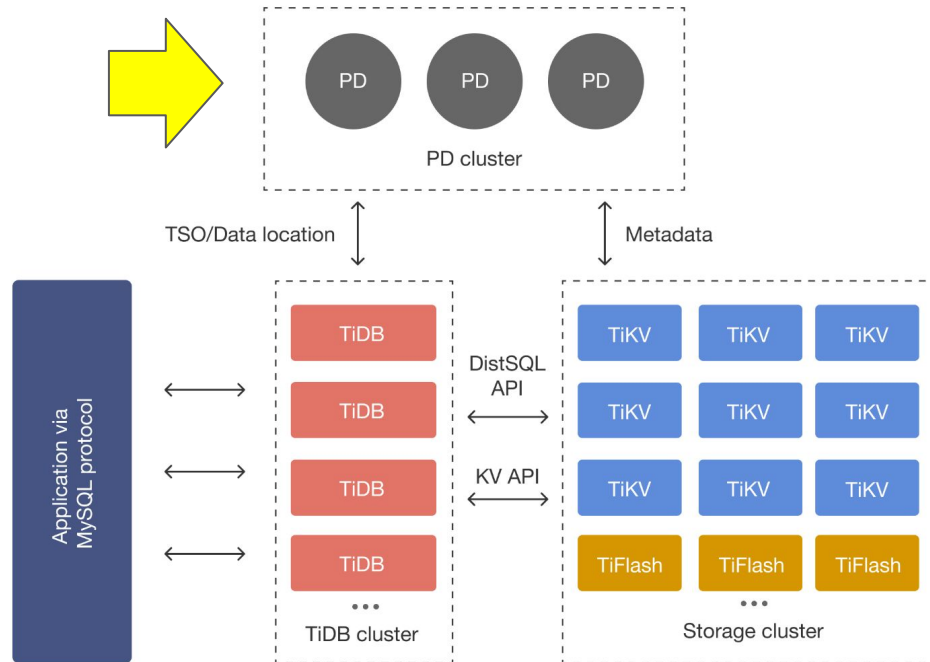
```
^C
# (demoblog uses TiDB by default, so no uri needed)
~/demoblog/demoblog
```

OK to stop the dm-master and dm-worker processes now, since traffic to MySQL is moved to TiDB.

Placement Driver (PD)

The placement driver:

- Has etcd embedded
- Is the Time Stamp Oracle (TSO): it gives out timestamps. These timestamps are used in transactions and in the MVCC system
- Takes care of data placement:
 - TiKV servers have labels
 - labels are used to ensure a raft group spans multiple availability zones
 - Splitting big data regions, merging small data regions, splitting hot data regions, evenly distributing regions across the cluster.



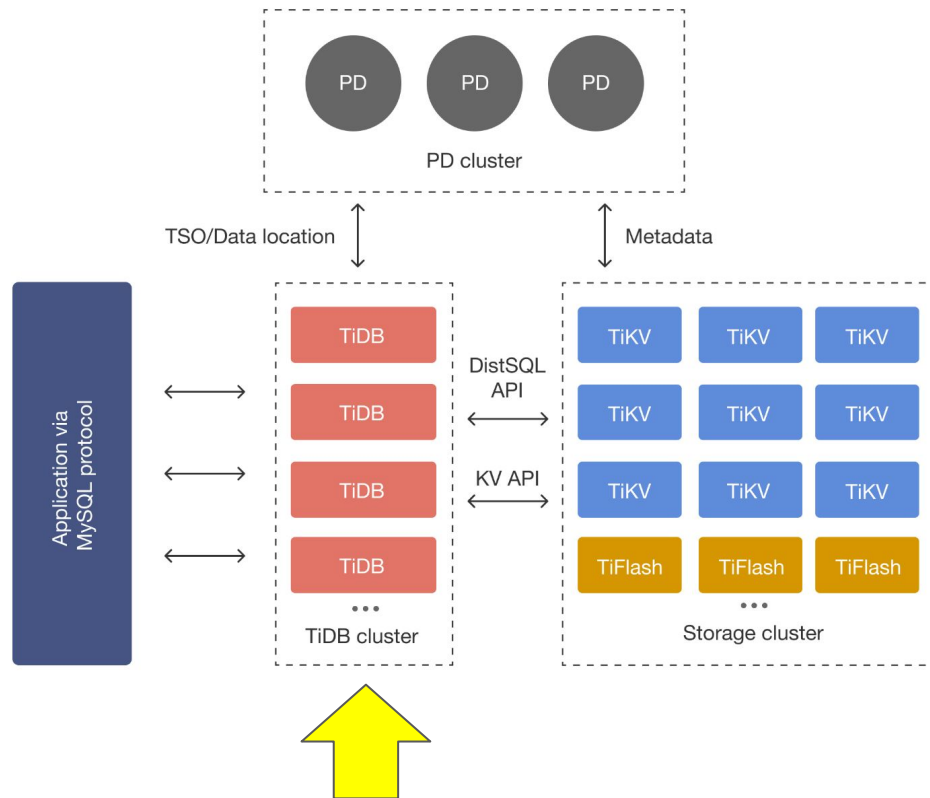
TiDB Server

TiDB Server is one of the components of the TiDB Platform. This is a bit confusing.

TiDB Server is written in Go and doesn't share any code with MySQL.

This implements the MySQL protocol, optimizer and executor.

TiDB sql node is stateless, it doesn't contain data.



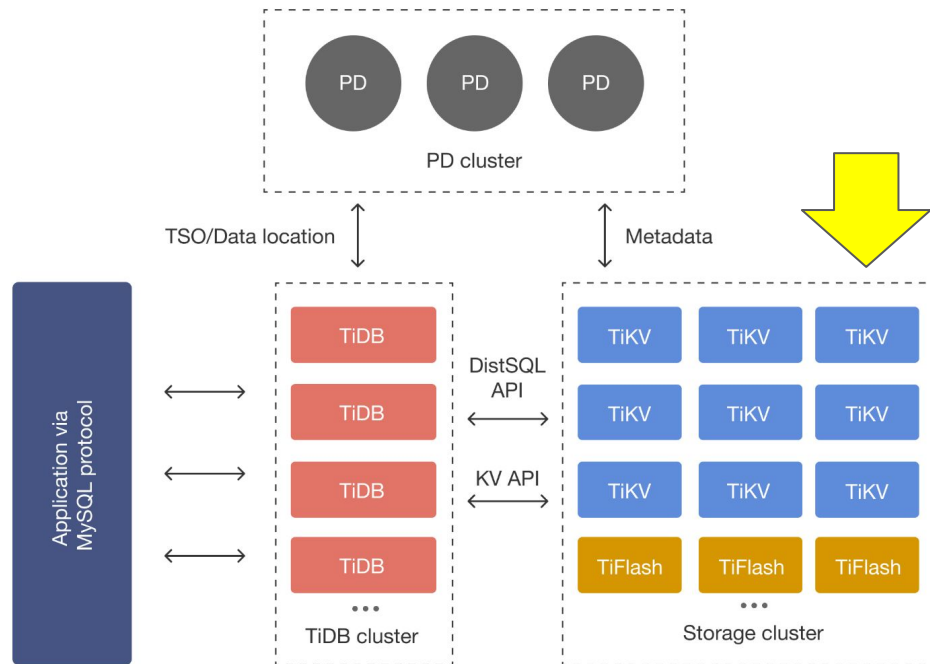
TiKV



TiKV is a distributed key-value store. This is a CNCF project.

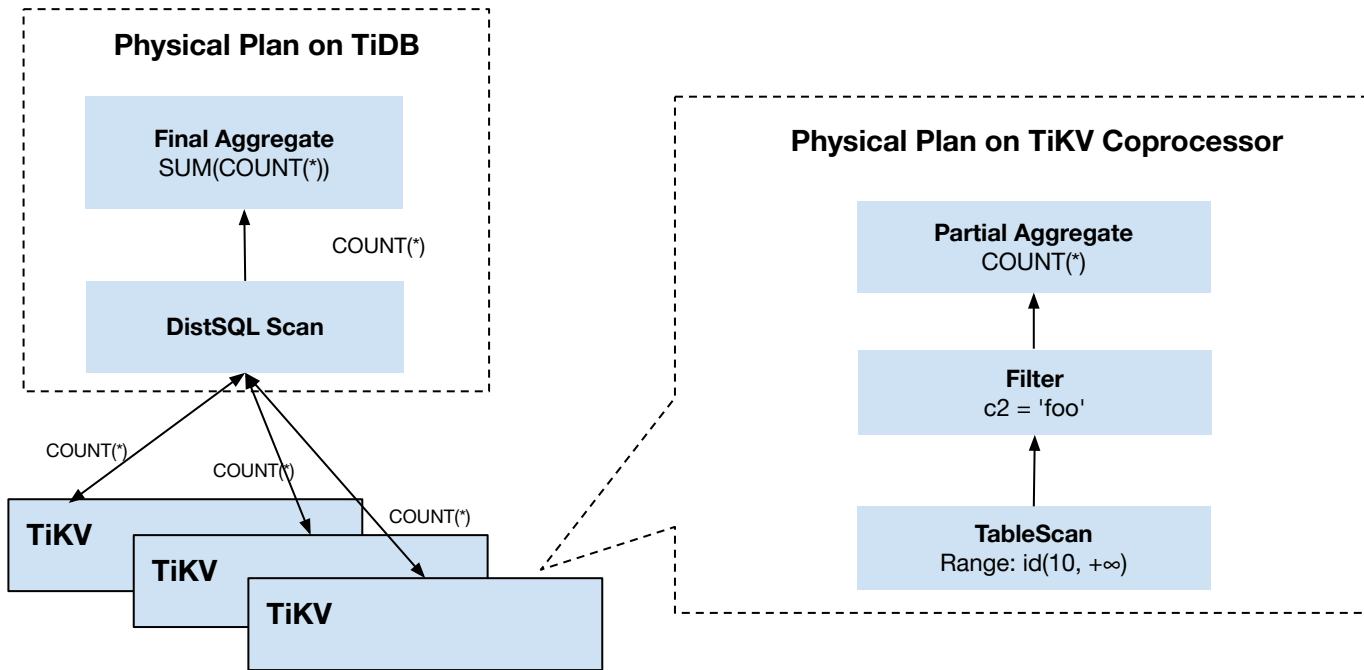
Database tables are stored with a RowID or PK as key and the columns as values.

A database table is split up into multiple data regions. Each data region is a raft group of (by default) three nodes.



Query Execution/Distributed Computing

```
SELECT COUNT(*) FROM t WHERE id > 10 AND c2 = 'foo';
```



TiDB Architecture

UPDATE orders
SET delivered = 1
WHERE id = 12345

Stateless SQL Layer

TiDB node 1

TiDB node 2

TiDB node 3

AZ 1

TiKV node 1

Region 5

Region 3

Region 4

TiKV node 4

Region 1

Region 6

Region 2

AZ 2

TiKV node 2

Region 1

Region 2

Region 3

TiKV node 5

Region 5

Region 6

Region 4

AZ 3

TiKV node 3

Region 2

Region 4

Region 5

TiKV node 6

Region 6

Region 1

Region 3

Scalability with TiDB

Both reads and writes can go to any of the TiDB nodes.

Scaling reads?

Add more nodes

Scaling writes?

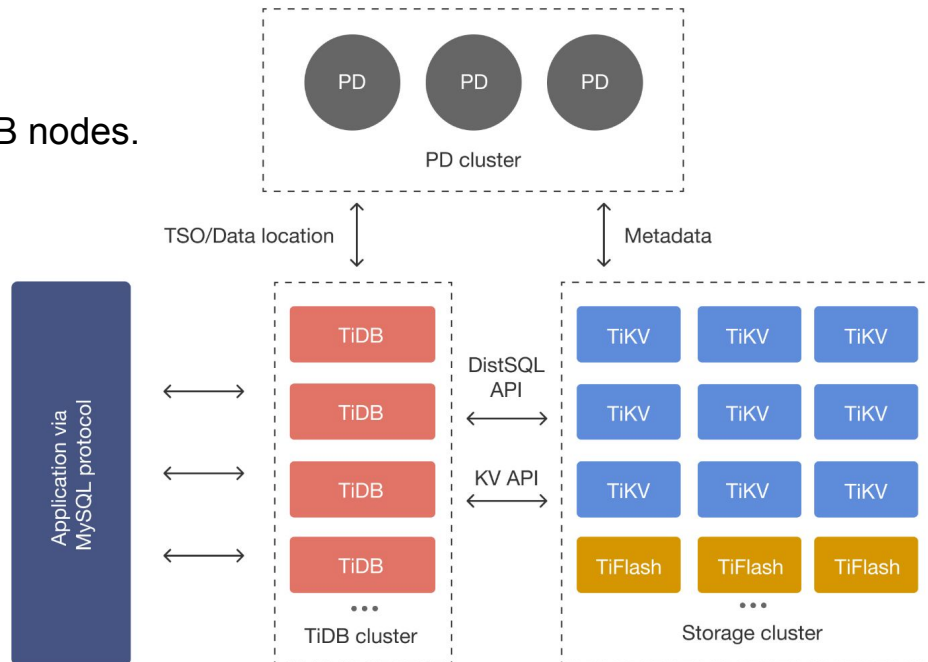
Add more nodes

Scaling data volume?

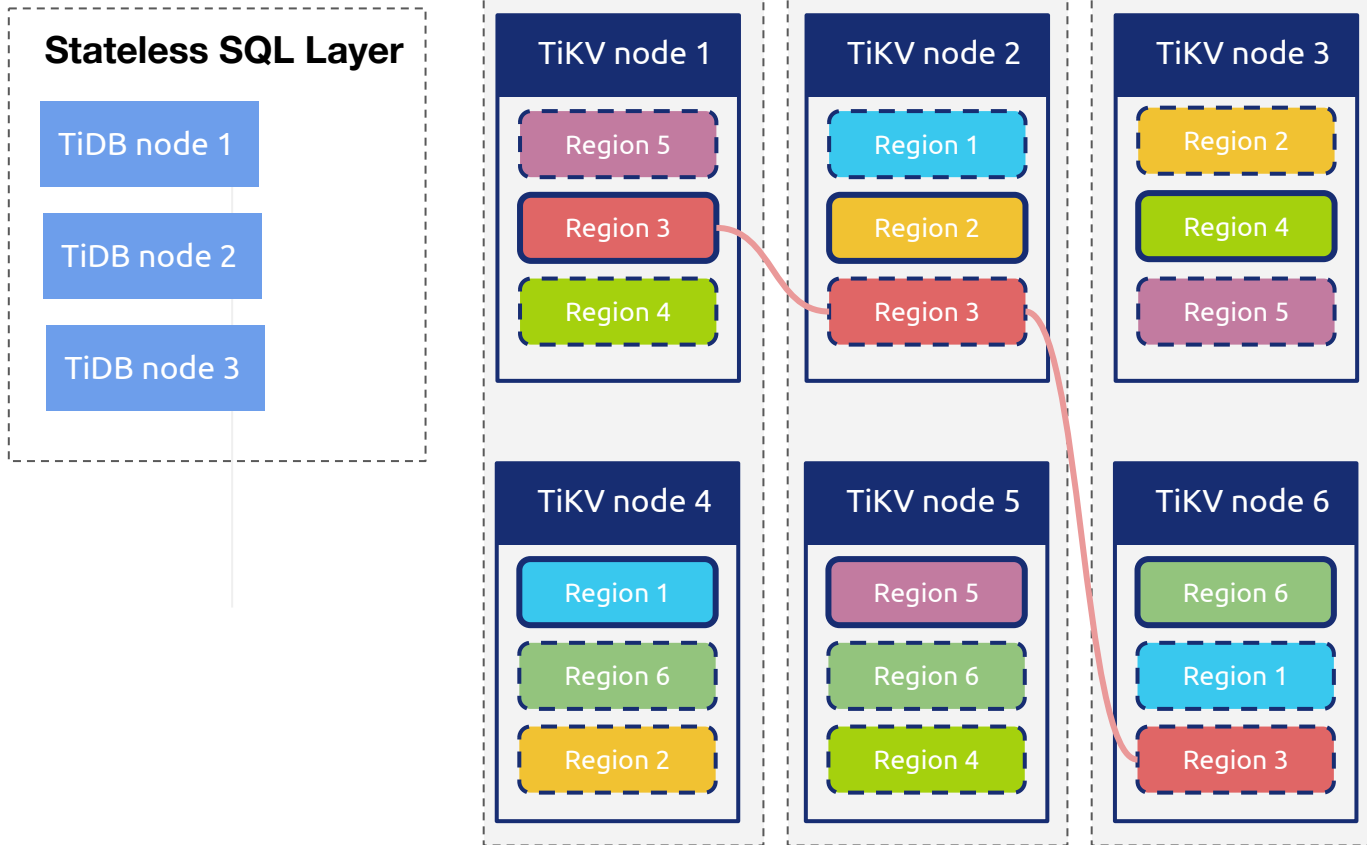
Add more nodes

Need to scale more?

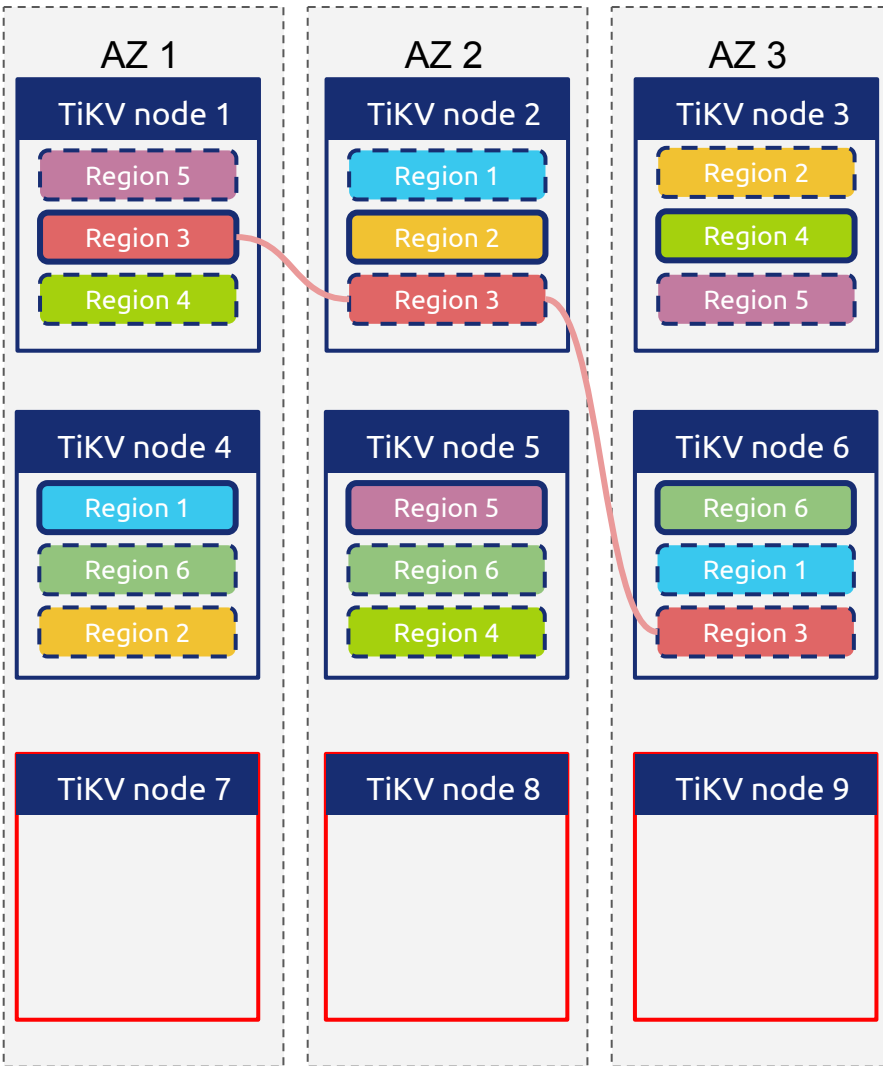
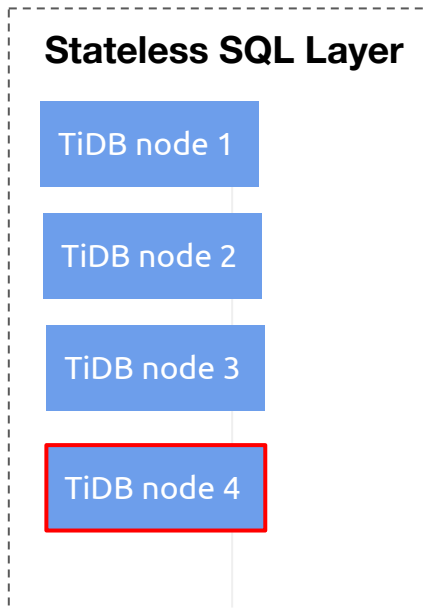
Add more nodes



TiDB Architecture



TiDB Architecture



TiDB Architecture

Stateless SQL Layer

TiDB node 1

TiDB node 2

TiDB node 3

TiDB node 4

AZ 1

TiKV node 1

Region 5

Region 3

TiKV node 4

Region 6

Region 2

TiKV node 7

Region 1

Region 4

AZ 2

TiKV node 2

Region 1

Region 3

TiKV node 5

Region 5

Region 4

TiKV node 8

Region 6

Region 2

AZ 3

TiKV node 3

Region 2

Region 4

TiKV node 6

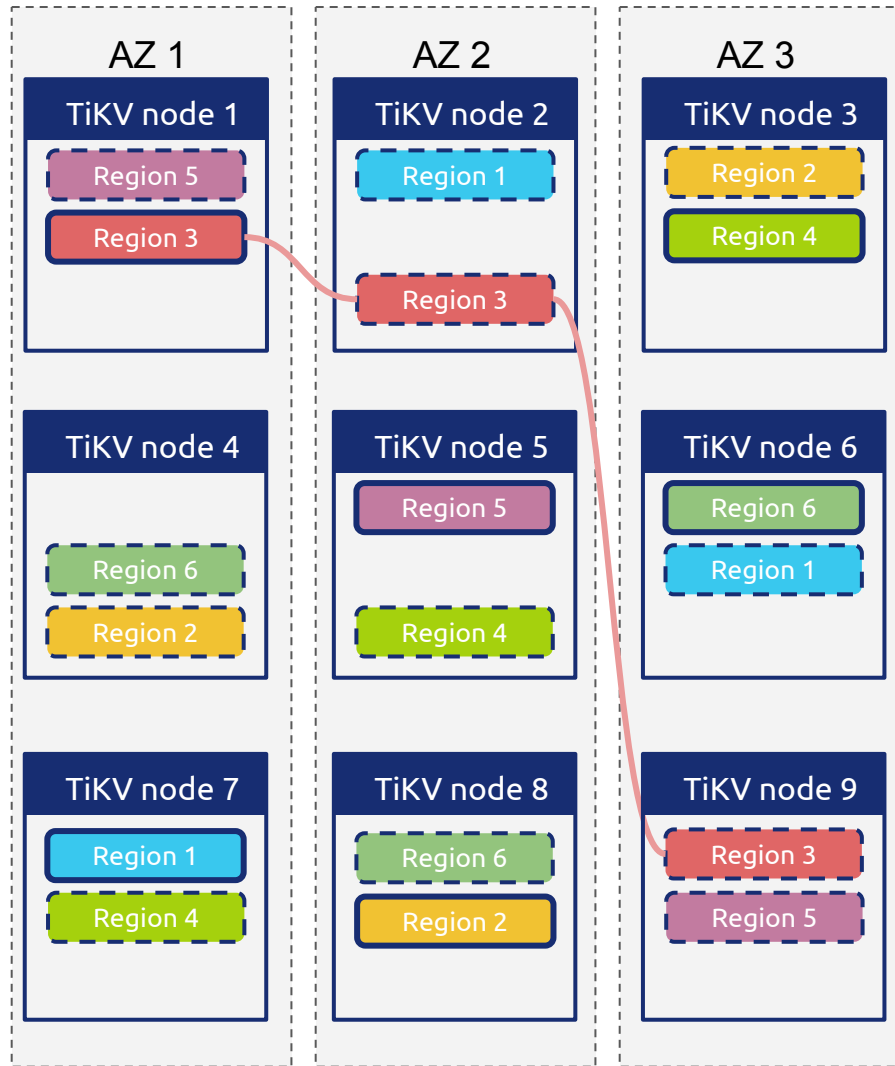
Region 6

Region 1

TiKV node 9

Region 3

Region 5



Lab excercises: TiKV

MySQL Client:

```
mysql -u root -h 127.0.0.1 -P 4000 blog
```

```
SHOW TABLE comments REGIONS;
```

Run loadGen (in a new ssh session)

```
cd demoblog/loadGen  
./loadGen
```

```
SHOW TABLE comments REGIONS;
```

```
SELECT REGION_ID, START_KEY, END_KEY, TABLE_NAME, INDEX_NAME FROM  
information_schema.TIKV_REGION_STATUS WHERE DB_NAME='blog';
```

Check Key Visualizer in TiDB Dashboard

```
ssh -L2379:127.0.0.1:2379 ...  
http://127.0.0.1:2379/dashboard/
```

Lab exercises

Scale out via TiUP Playground:

```
tiup playground display
tiup playground scale-out --kv 2
tiup playground display
tiup ctl:v7.2.0 pd store
curl http://127.0.0.1:2379/pd/api/v1/stores
```

SQL:

```
SELECT * FROM information_schema.TIKV_STORE_STATUS;
```

```
SELECT REGION_ID, START_KEY, END_KEY, TABLE_NAME, INDEX_NAME FROM
information_schema.TIKV_REGION_STATUS WHERE DB_NAME='blog';
```

```
SELECT * FROM information_schema.TIKV_REGION_PEERS WHERE
REGION_ID=<id from one of the regions above>;
```

```
mysql> SELECT REGION_ID,START_KEY,END_KEY,TABLE_NAME,INDEX_NAME FROM information_schema.TIKV_REGION_STATUS WHERE DB_NAME='blog';
+-----+-----+-----+-----+-----+
| REGION_ID | START_KEY | END_KEY | TABLE_NAME | INDEX_NAME |
+-----+-----+-----+-----+-----+
| 131 | 7480000000000000FF72000000000000F8 | 7480000000000000FF75000000000000F8 | posts | NULL |
| 133 | 7480000000000000FF75000000000000F8 | 7480000000000000FF77000000000000F8 | comments | post_id |
| 133 | 7480000000000000FF75000000000000F8 | 7480000000000000FF77000000000000F8 | comments | NULL |
| 14 | 7480000000000000FF77000000000000F8 | 748000FFFFFFFFFFFFFFFF90000000000000F8 | authors | NULL |
| 129 | 7480000000000000FF71000000000000F8 | 7480000000000000FF72000000000000F8 | ticker | NULL |
+-----+-----+-----+-----+-----+
5 rows in set (0.01 sec)
```

```
mysql> SELECT * FROM information_schema.TIKV_REGION_PEERS WHERE REGION_ID=133;
+-----+-----+-----+-----+-----+-----+-----+
| REGION_ID | PEER_ID | STORE_ID | IS_LEARNER | IS_LEADER | STATUS | DOWN_SECONDS |
+-----+-----+-----+-----+-----+-----+-----+
| 133 | 134 | 1 | 0 | 1 | NORMAL | NULL |
| 133 | 193 | 136 | 0 | 0 | NORMAL | NULL |
| 133 | 269 | 135 | 0 | 0 | NORMAL | NULL |
+-----+-----+-----+-----+-----+-----+-----+
3 rows in set (0.00 sec)
```



```
mysql> SELECT REGION_ID, TABLE_ID, INDEX_NAME, TIDB_DECODE_KEY(START_KEY) 'start', TIDB_DECODE_KEY(END_KEY) 'end', APPROXIMATE_SIZE
-> FROM information_schema.TIKV_REGION_STATUS WHERE DB_NAME='blog' AND TABLE_NAME='comments'
-> ORDER BY INDEX_NAME, START_KEY, END_KEY;
```

REGION_ID	TABLE_ID	INDEX_NAME	start	end	APPROXIMATE_SIZE
149	100	NULL	7480000000000000FF645F720000000000FA	{"id":1357949,"table_id":"100"}	116
134	100	NULL	{"id":1357949,"table_id":"100"}	{"table_id":101}	131
123	100	post_id	{"table_id":100}	{"index_id":1,"index_vals":{"post_id":"1"},"table_id":100}	1
125	100	post_id	{"index_id":1,"index_vals":{"post_id":"1"},"table_id":100}	{"index_id":1,"index_vals":{"post_id":"1"},"table_id":100}	1
152	100	post_id	{"index_id":1,"index_vals":{"post_id":"1"},"table_id":100}	{"index_id":1,"index_vals":{"post_id":"5"},"table_id":100}	91
137	100	post_id	{"index_id":1,"index_vals":{"post_id":"5"},"table_id":100}	7480000000000000FF645F698000000000FF0000020003800000FF0000000001000000FC	64

6 rows in set, 2 warnings (0.00 sec)

- TIDB_DECODE_KEY() to decode keys

Lab excercises

Add yet another TiKV node, so we safely can remove one:

```
tiup playground scale-out --kv 1
```

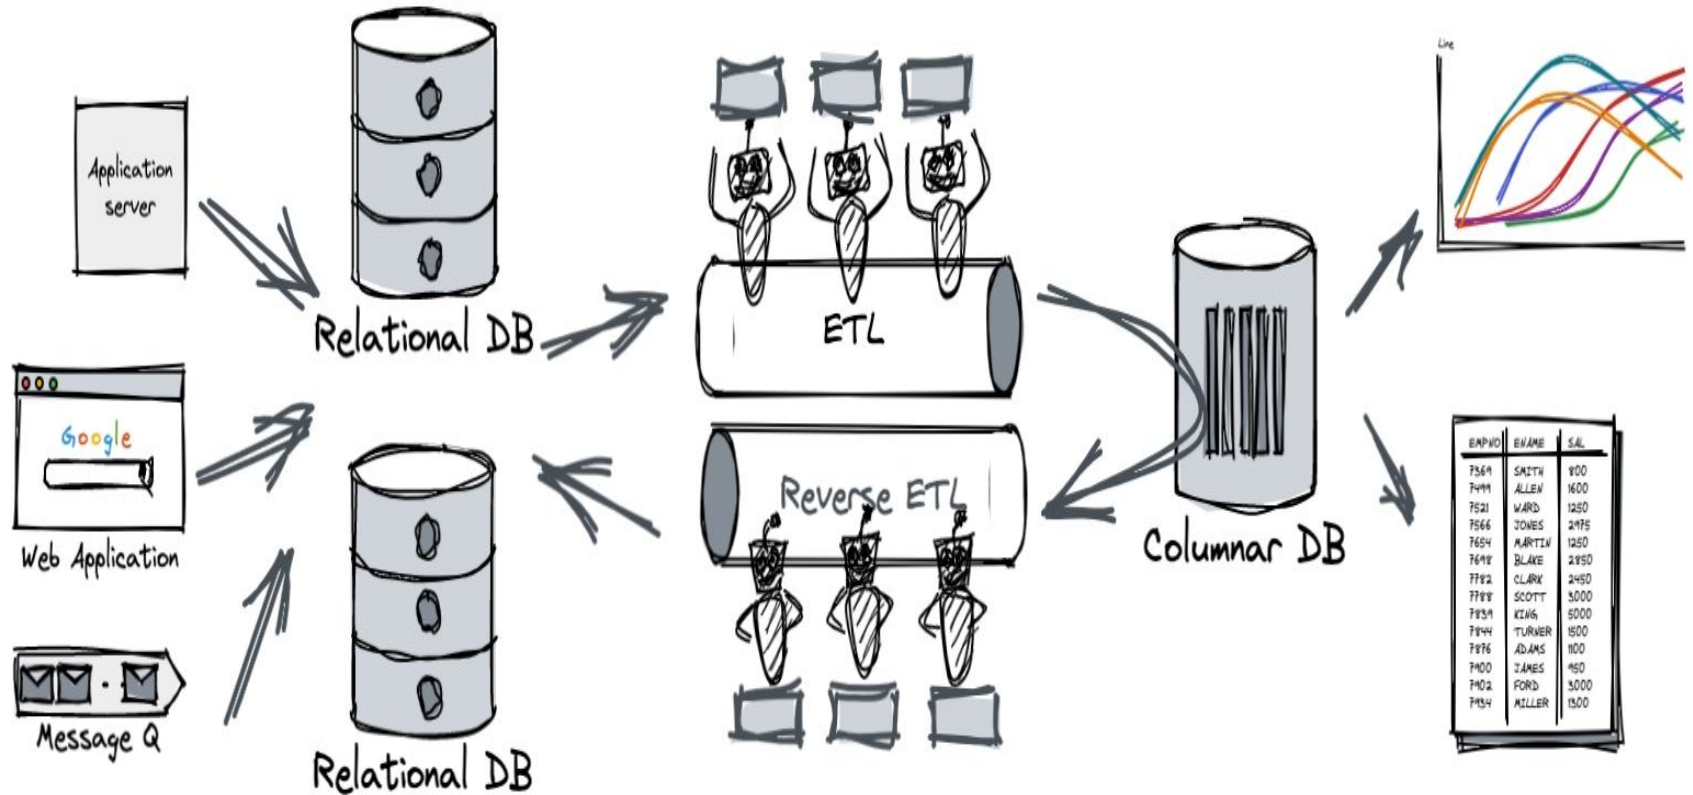
Scale in via TiUP Playground:

```
tiup playground display
```

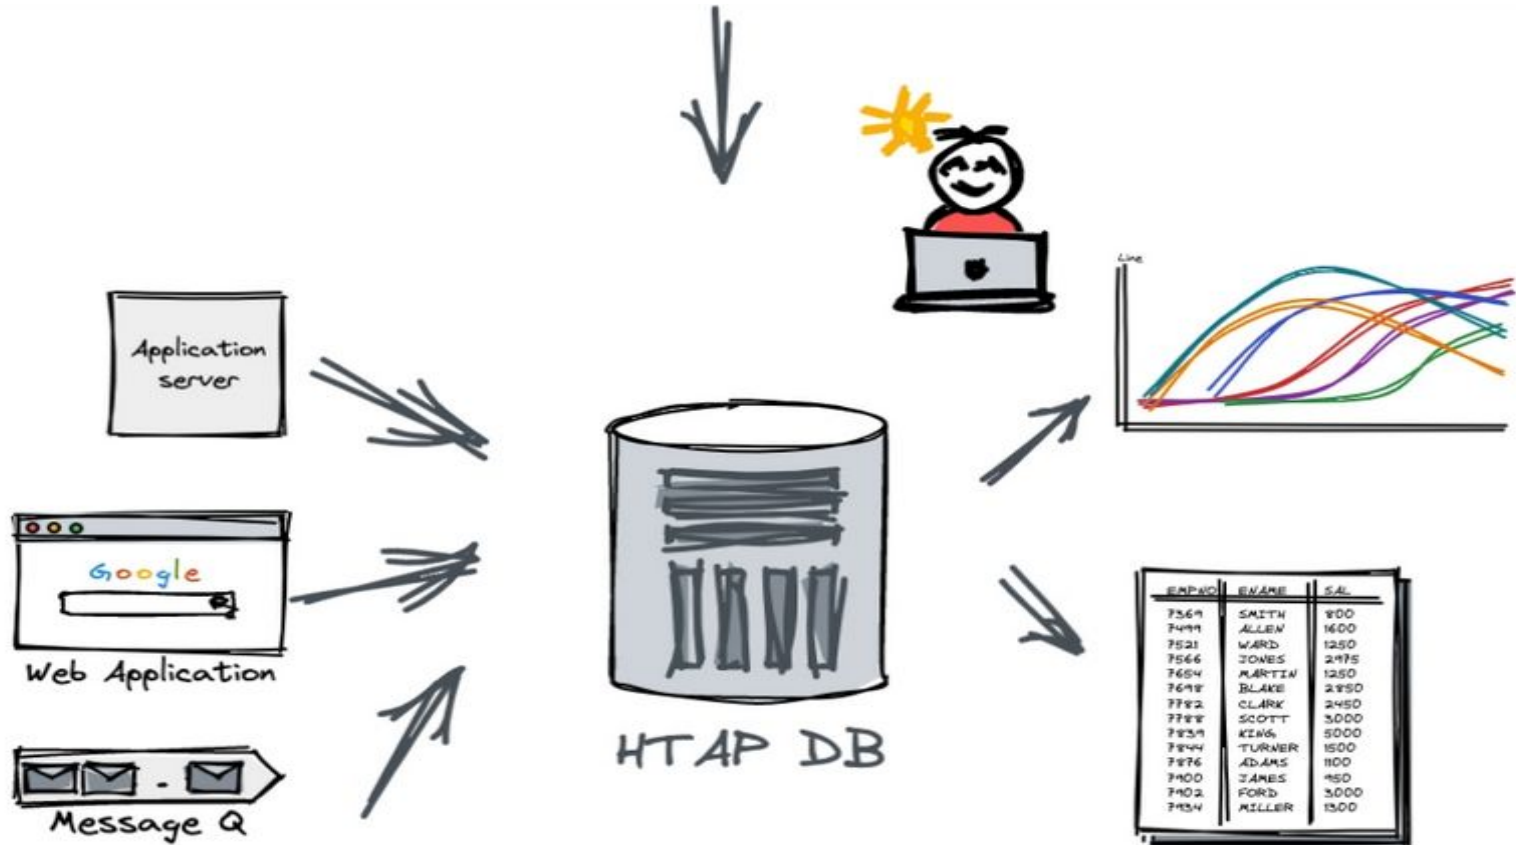
```
tiup playground scale-in --pid $pid
```

Scaling in TiDB nodes will be quick. TiKV might need to move data.

Anyone recognizes similar architecture?



What if it could be simplified to this?



TiDB Architecture with TiFlash

UPDATE orders
SET delivered = 1
WHERE id = 12345

Stateless SQL Layer

TiDB node 1

TiDB node 2

TiDB node 3

TiDB node 4

AZ 1

TiKV node 1

Region 5

Region 3

Region 4

TiKV node 4

Region 1

Region 6

Region 2

TiFlash 1

Region 1

Region 2

Region 6

AZ 2

TiKV node 2

Region 1

Region 2

Region 3

TiKV node 5

Region 5

Region 6

Region 4

TiFlash 2

Region 3

Region 6

AZ 3

TiKV node 3

Region 2

Region 4

Region 5

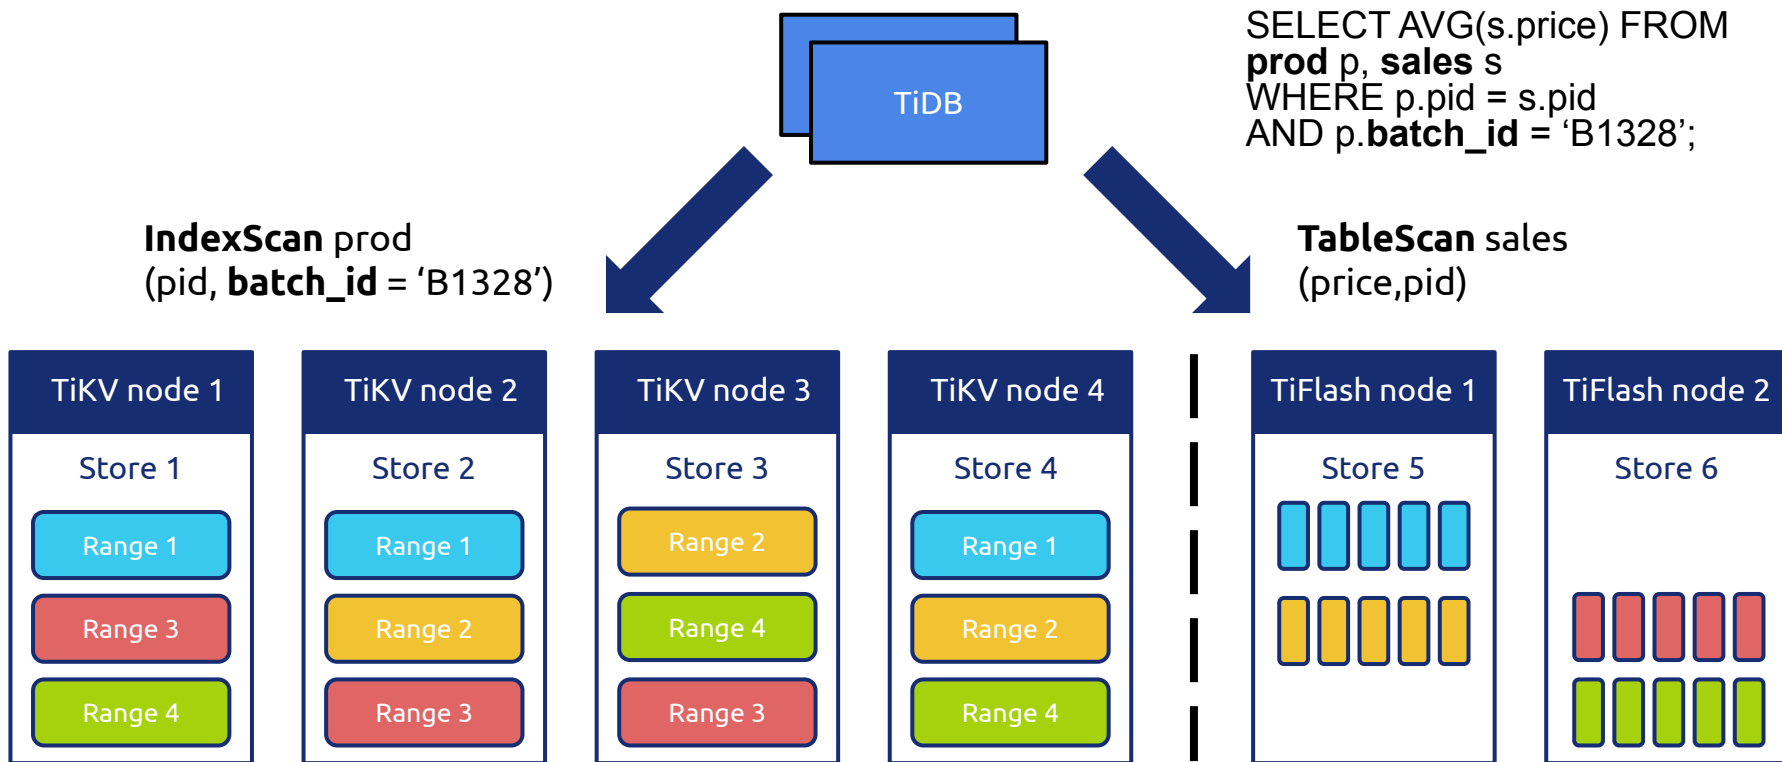
TiKV node 6

Region 6

Region 1

Region 3

The optimizer can pick the best store for each case



Lab excercises: TiFlash

Setup TiFlash replicas

```
ALTER TABLE posts SET TIFLASH REPLICA 1;  
ALTER TABLE comments SET TIFLASH REPLICA 1;
```

```
EXPLAIN SELECT p.title, SUM(c.likes) FROM posts p LEFT JOIN comments c ON  
c.post_id=p.id GROUP BY p.id;
```

Run loadGen for a bit and run explain again.

Try `EXPLAIN ANALYZE <query>`

Try engine isolation:

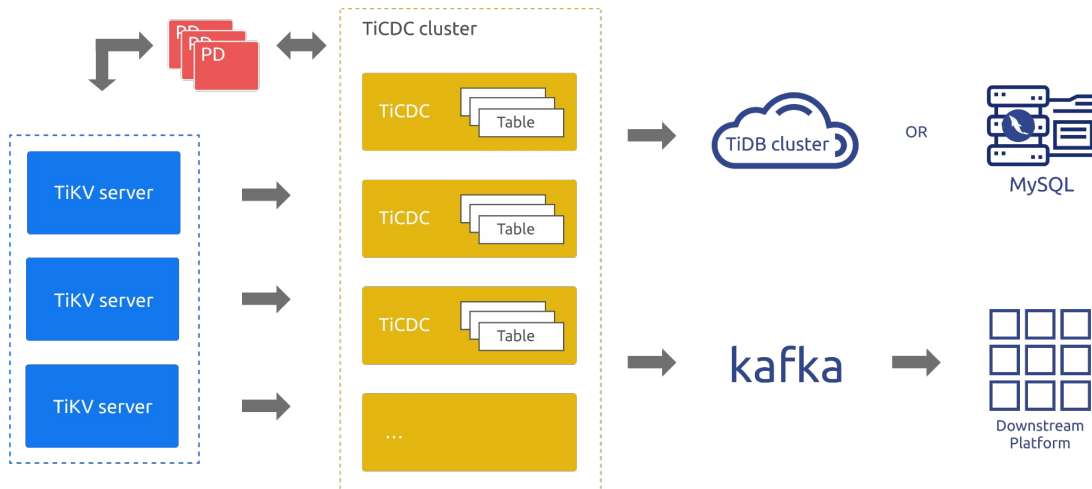
```
SET SESSION tidb_isolation_read_engines = "tikv";  
EXPLAIN ...  
SET SESSION tidb_isolation_read_engines = "tiflash,tikv";
```

Only adding TiFlash for comments allows joins between comments and posts. This combines TiKV and TiFlash

Batteries included: TiCDC

TiCDC

- Change Data Capture
- Send events to Kafka, MySQL or another TiDB Cluster.
- Also high available



Lab excercises: TiCDC

Try ticketStat:

- Setup Kafka
- Setup TiCDC changefeed to kafka
- Run tickerStat (consumes kafka)

```
cd demoblog/tickerStat
```

```
cat README.md
```

Compatibility

TiDB is compatible with the MySQL protocol and syntax of MySQL 5.7 and MySQL 8.x. There are a few differences.

Row size is limited to 120 MiB.

TiDB differences:

- AUTO_INCREMENT, AUTO_RANDOM (but there is a MySQL compatible mode)
- Explain output is different
- No triggers, stored procedures, etc: Use TiCDC changefeed instead.
- No fulltext index
- No geospatial support

<https://docs.pingcap.com/tidb/stable/mysql-compatibility>

Lab excercises: Compatibility

Try queries:

```
SELECT p.title,p.likes,SUM(c.likes) 'liked comments'  
FROM posts p  
LEFT JOIN comments c ON c.post_id=p.id GROUP BY p.id;
```

Try some tools like MySQL Workbench, DBeaver, etc. This might need portforwarding of port 4000: `ssh -L4000:127.0.0.1:4000 ...`

<https://docs.pingcap.com/tidb/stable/mysql-compatibility>

Most ORM's just work with MySQL compatibility. There are some that implement TiDB specific features and settings (e.g. default port as 4000)

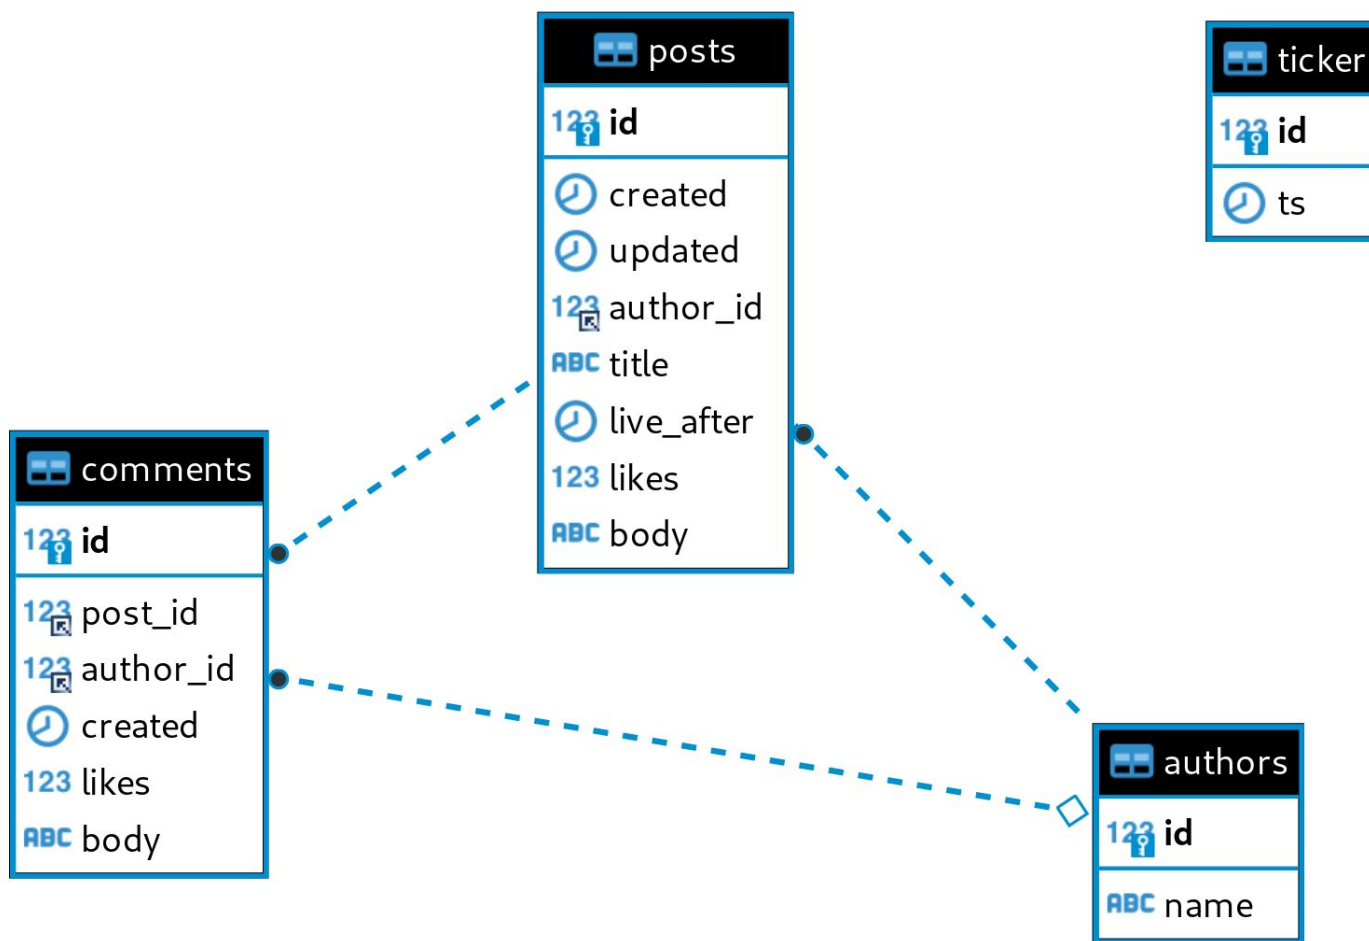
Lab excercises: Online DDL

Drop and Add INDEX

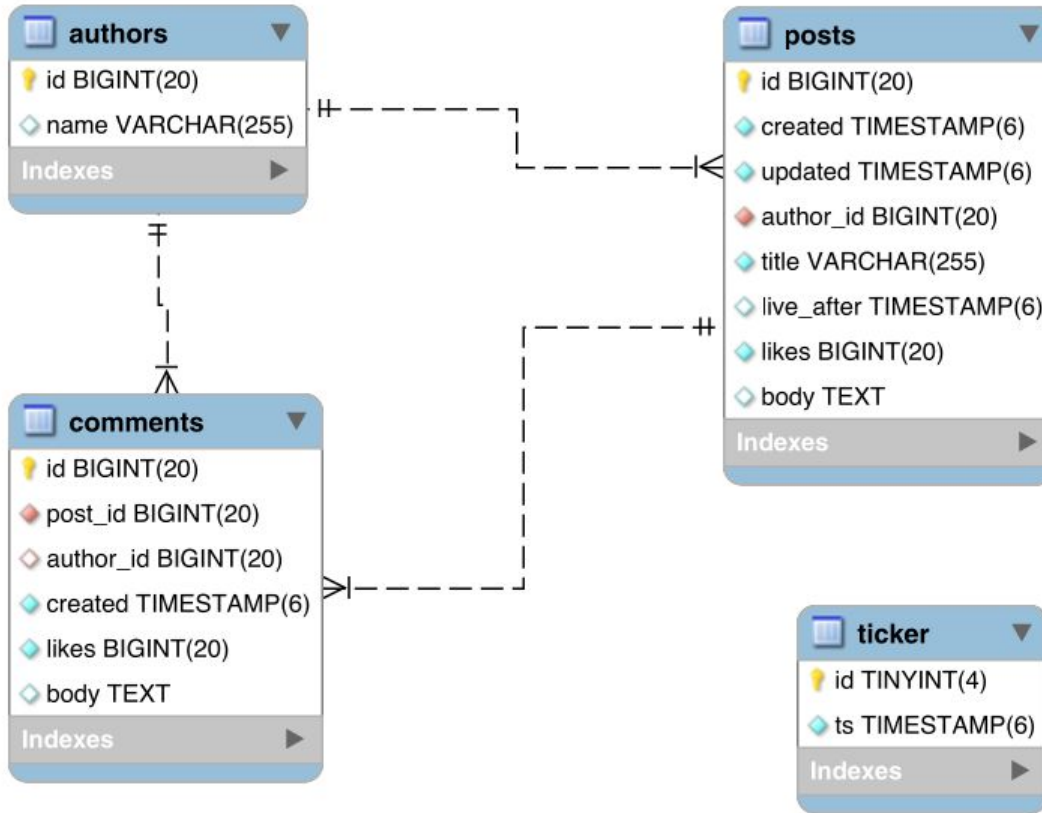
```
ALTER TABLE comments DROP INDEX post_id;  
ALTER TABLE comments ADD INDEX post_id (post_id);
```

Foreign Keys

```
ALTER TABLE comments  
    ADD CONSTRAINT fk_postid FOREIGN KEY (post_id) REFERENCES posts(id),  
    ADD CONSTRAINT fk_authorid FOREIGN KEY (author_id) REFERENCES authors(id);  
  
ALTER TABLE posts  
    ADD CONSTRAINT fk_authorid FOREIGN KEY (author_id) REFERENCES authors(id);  
  
-- This should fail as there is no author with that ID.  
UPDATE posts SET author_id=2 WHERE id=1;
```



DBeaver



Time travel with TiDB

- Data is never actually removed directly, it is garbage collected instead.
- This allows you to see any version between now and the point of the GC.
- Use `SELECT ... FROM ... AS OF TIMESTAMP`
- Works for both table data and table structure
- Dumping can dump data of a specific time
- Default GC life time is 10 minutes
- `FLASHBACK` statement can recover the last version of a table, database or a full cluster

Lab exercises

- Time travel: AS OF...

```
SELECT post_id,COUNT(*),SUM(likes)
FROM comments AS OF TIMESTAMP '2023-07-28 15:25:11'
GROUP BY post_id;
```

- FLASHBACK TABLE & BR

This was in dev environment right?!

```
DROP TABLE comments;
```

Oh no, prod :o (loadGen probably crashed...)

```
FLASHBACK TABLE comments;
```

Restart the loadGen and it should work again \o/

Lab excercises

Now check the various dashboards:

- <http://127.0.0.1:2379/dashboard> PD dashboard (root/)
- <http://127.0.0.1:3000/> Grafana (admin/admin)

Backup and Restore

BR: Backup and Restore

Every TiKV node writes its own backup to S3 or other shared storage. Similar for restores. This makes backups scalable.

Lab excercises: BR

Create a backup

```
tiup br backup full -s file:///tmp/b1
```

Decode backup metadata

```
tiup br validate decode -s /tmp/b1/
```

Get the TSO of the backup

```
jq .end_version /tmp/b1/backupmeta.json
```

Get the timestamp for the TSO

```
tiup ctl:v7.2.0 pd tso $tso
```

Similar, via SQL: `SELECT TIDB_PARSE_TSO($tso)`

Lab excercises: BR & TSO

```
WITH tso AS (  
    SELECT 443165077170552833 v  
)  
SELECT  
    v 'TSO raw',  
    (v >> 18)/1000 'TSO unix timestamp',  
    FROM_UNIXTIME((v >> 18) / 1000) 'TSO timestamp',  
    TIDB_PARSE_TSO(v),  
    v - ((v >> 18) << 18) 'TSO logical part'  
FROM tso\G
```

443165077170552833 in binary:

0000011000100110011100000100101000101001100111 00000000000000000000

1

< 46 bits physical timestamp >< 18 bit logical

>

Dumpling and Lightning

Dumpling does parallel export of data from TiDB or MySQL

Lightning does:

- Data loading directly into TiKV (ingesting SST files, "local" backend)
- Data loading via TiDB (via SQL, "tidb" backend)

By default filenames are used as metadata, just like MyDumper and MyLoader.

Files can be local or on S3. (Not on AWS? use S3 compatible storage like MinIO)

TiDB Cloud has a webinterface for importing data from S3.

Lab exercises: Dumping

```
tiup dumping --database blog --filetype csv -o /tmp/dump  
find /tmp/dump
```

Lab excercises: Lightning

Create data:

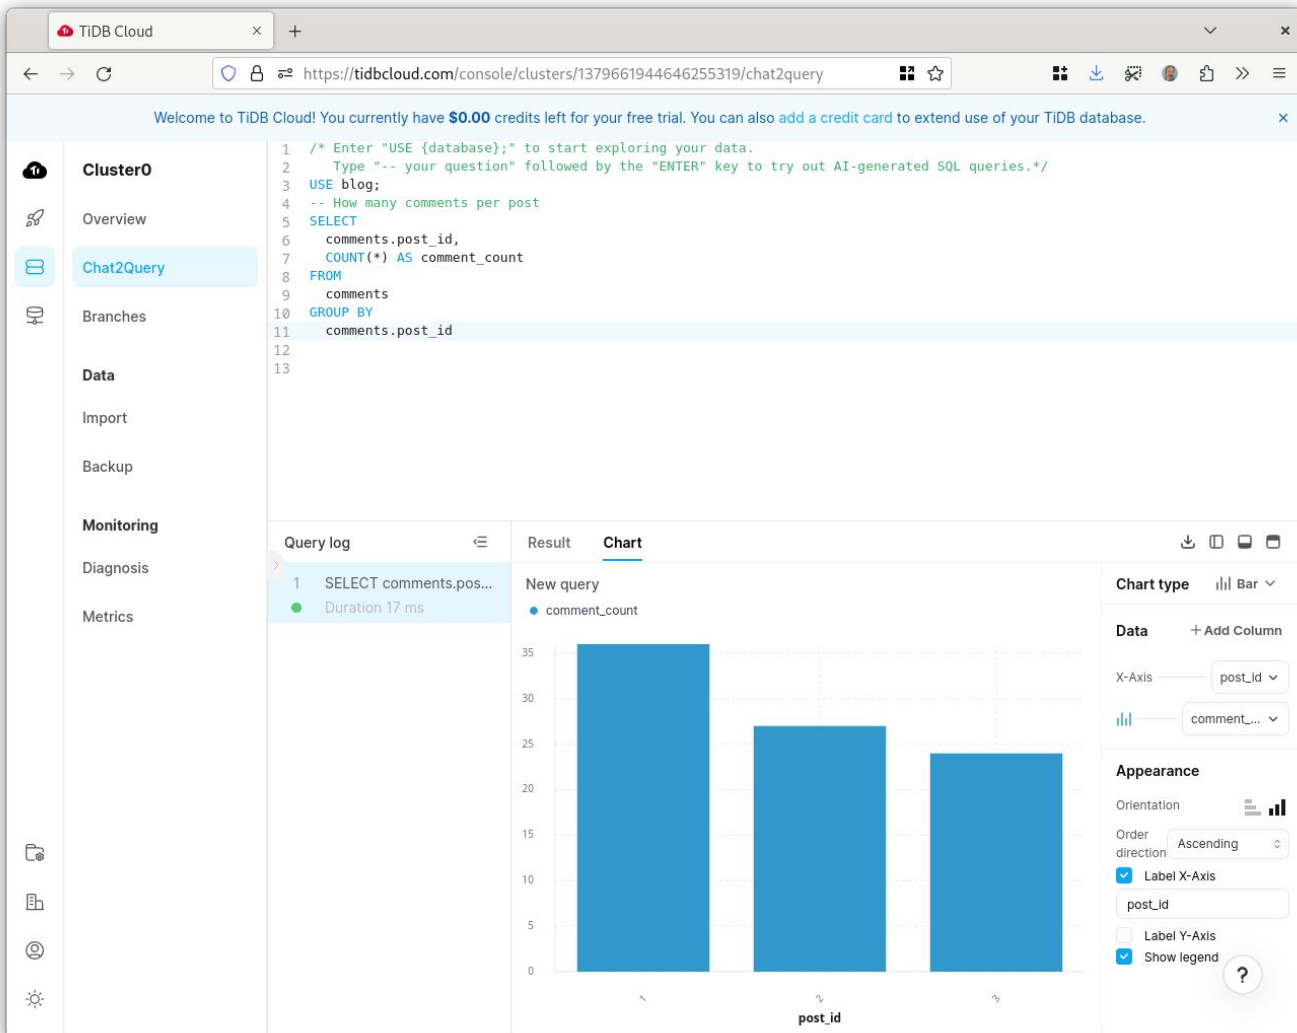
```
mkdir /tmp/authors
cat > /tmp/authors/blog.authors.000000000.csv
"id","name"
2,"Jane Doe"
```

Import data:

```
tiup tidb-lightning -d /tmp/authors -backend tidb -tidb-port
4000
```

Check via SQL:

```
TABLE authors;
```



Chat2Query:

- AI Generated SQL
- Charts

TiDB Cloud

https://tidbcloud.com/console/data-service

Welcome to TiDB Cloud! You currently have **\$0.00** credits left for your free trial. You can also add a [credit card](#) to extend use of your TiDB database.

Data Service BETA

search

foo

GET /posts

GET /posts

Cluster0 Run

```
1 /* Getting Started:
2  Enter "USE {database};" before entering your SQL statements.
3  Type "--your question" + Enter to try out AI-generated SQL queries
4  Declare a parameter like "Where id = ${arg}".
5  */
6 USE blog;
7 SELECT JSON_ARRAYAGG(id) FROM posts;
8
```

Result HTTP Response

```
{
  "root": {
    "type": "sql_endpoint",
    "data": {
      "columns": [
        {
          "col": "JSON_ARRAYAGG(id)",
          "data_type": "JSON",
          "nullable": true,
          "rows": [
            {
          "JSON_ARRAYAGG(id)": [1, 2, 3]
        }
      ]
    }
  },
  "result": {
    "code": 200,
    "message": "Query OK!",
    "start_ms": 1690387046659,
    "end_ms": 1690387046674,
    "latency": "15ms",
    "row_count": 1,
    "row_affect": 0,
    "limit": 50
  }
}
```

Properties Params Schema

Basic

Path

/posts/ids

Endpoint URL [Show Code Example](#)

https://eu-central-1.data.tidbcloud.com/api/v1beta/app/dataapi-ErkrdrGxa/endpoint/posts/ids

Request Method GET

Timeout(ms) 5000

Max Rows 50

Description

Input description

Advanced

Tag

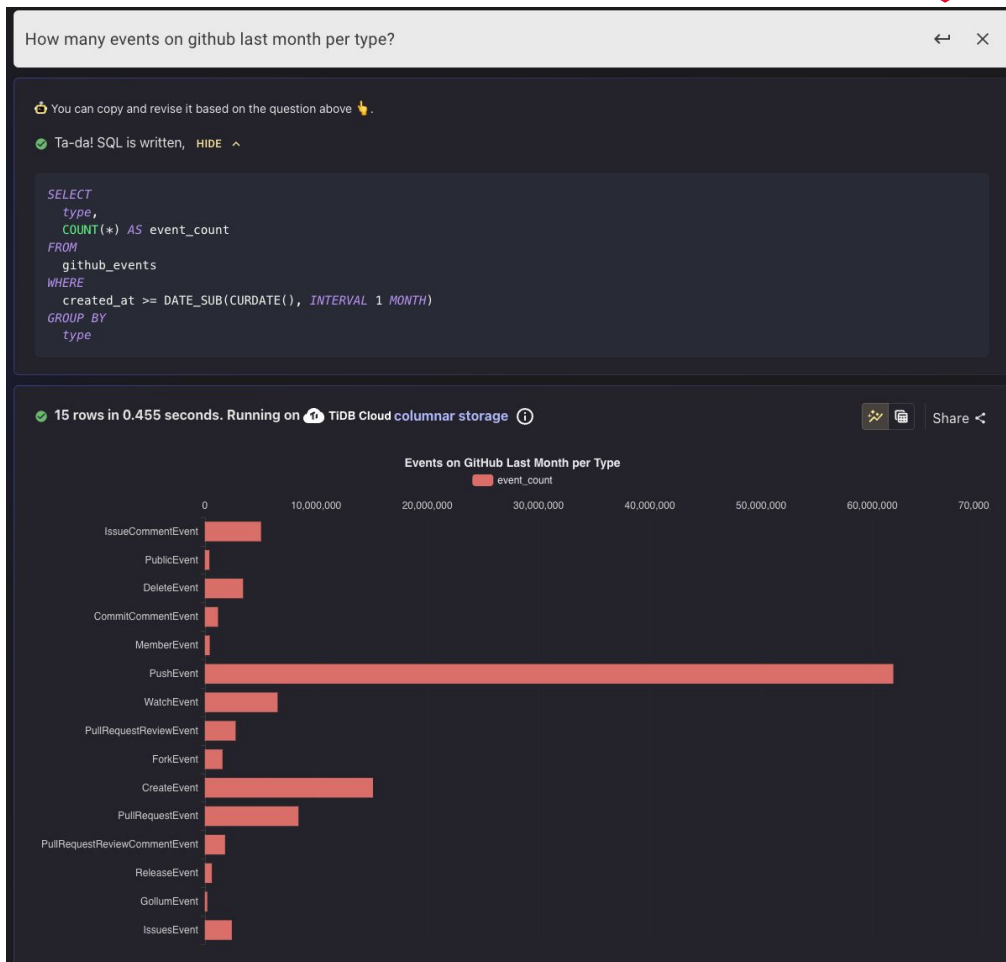
Request quotas are rate-limited at 100 requests per minute (rpm) per API Key. For an increase in quota, please [contact](#).



Data Service:
HTTP REST interface
based on SQL queries

OSS Insight

- <https://ossinsight.io>
- Built on TiDB.
- 6 Billion rows of GitHub events (Push, Pull, Watch/Star, Issue etc.)
- Allows you to look at the SQL, including explain plans.
- No ETL.
- Chat2Query



A real-world TiDB cluster in production

Tweets from Xiaoguang Sun, Head of Infrastructure @ Zhihu



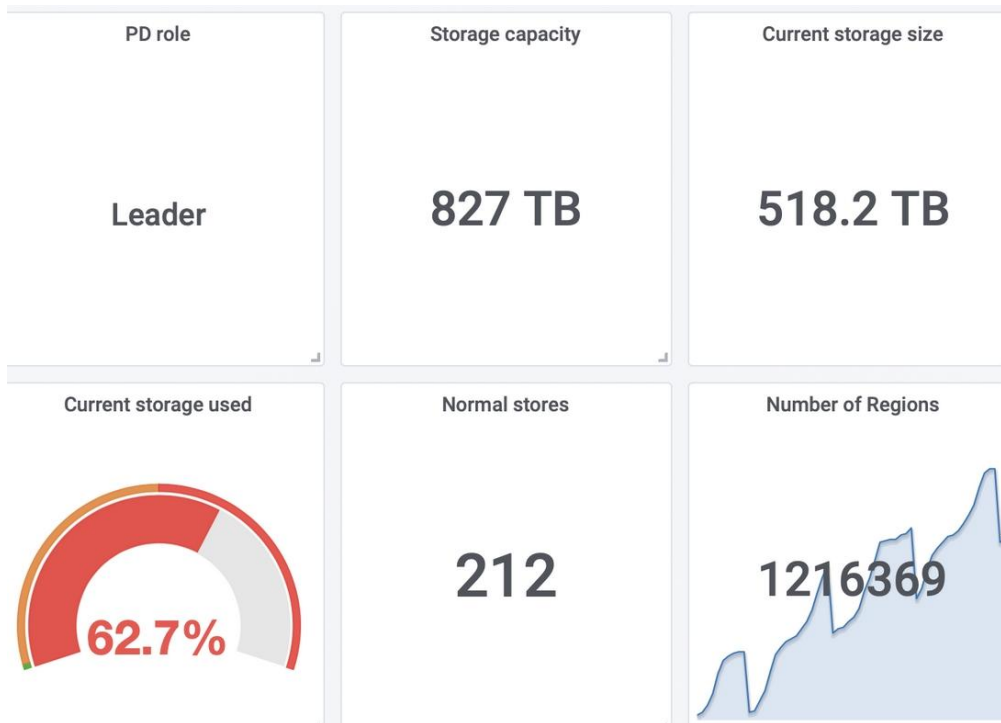
Xiaoguang Sun
@xgsun

This is perhaps the largest single TiDB cluster in the world: 168 TiKV instances / 21 TiDB instances / 1,820 billion rows of records / 318TB of data / peak read 100 million rows per second / peak write 87,000 rows per second.

It wouldn't have been possible if TiDB hadn't been there

[Translate Tweet](#)

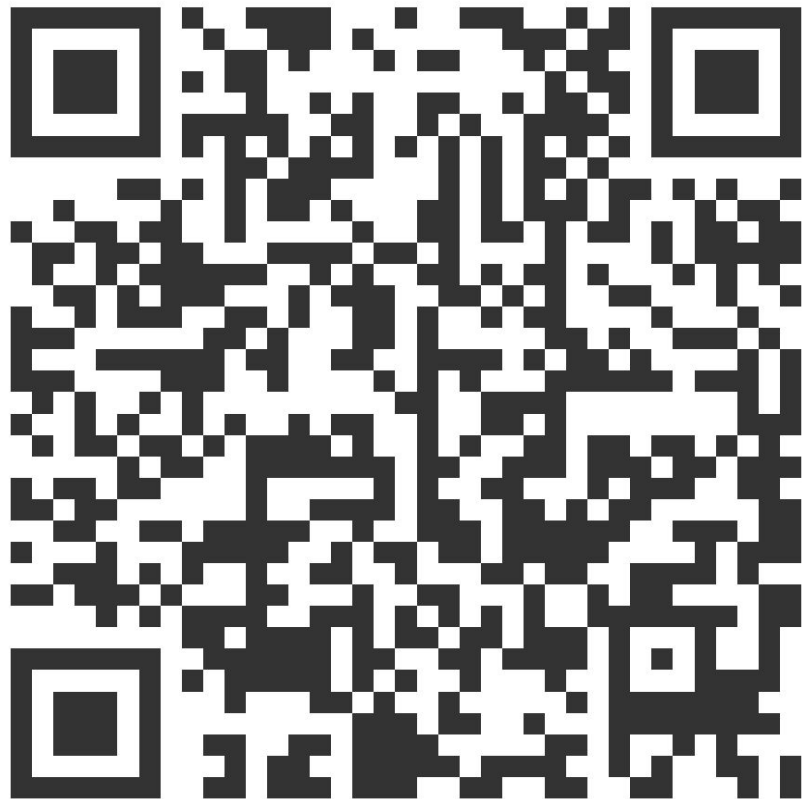
Only two engineers are maintaining
TiDB at Zhihu →



*Dashboard Screenshot of the biggest
TiDB cluster running in Zhihu.com*

Conclusions and next steps

- TiDB is MySQL compatible
- TiDB has fully automatic scaling
- TiDB has high availability builtin
- TiDB is developer friendly
- TiDB allows for online DDL
- TiDB can be deployed
 - on-prem with TiUP Cluster
 - on-prem via Kubernetes with TiDB Operator
 - via TiDB Cloud on AWS or GCP
- Ask questions via Slack (via QR code)
- Give TiDB a try for your application



High availability with MySQL

An often-used setup with MySQL is to have a primary and multiple replicas.
Then if the primary fails you promote a replica.

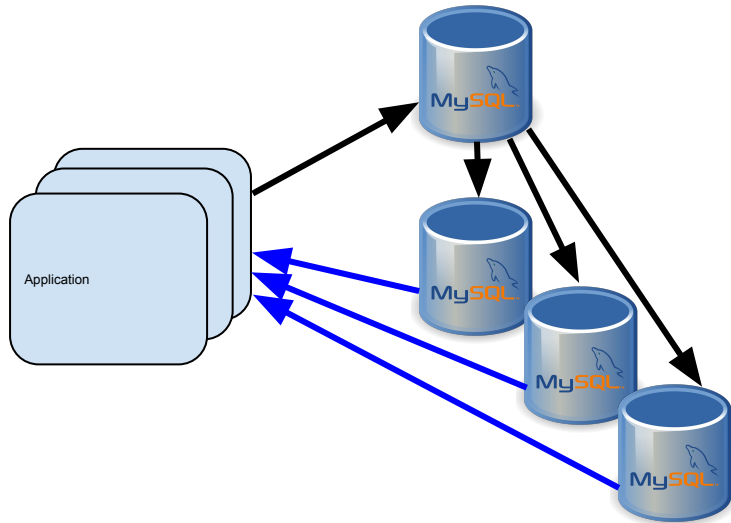
Replica promotion is not automated.

Is your replica up to date?

Use a load-balancer for service discovery?

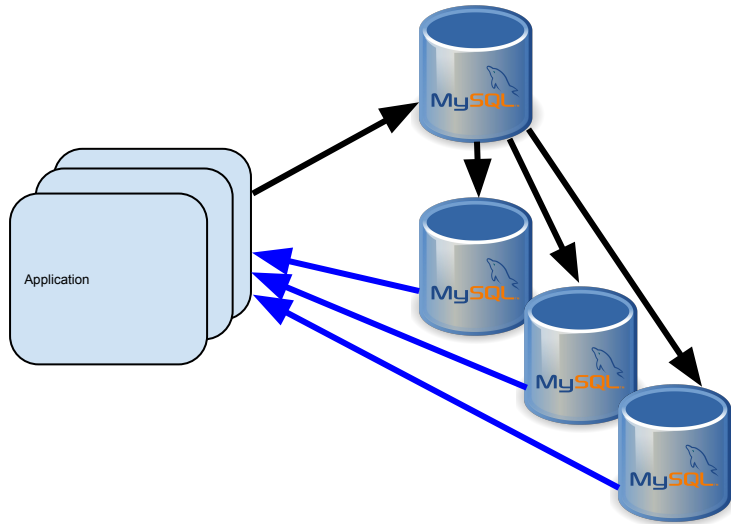
How to fail back once the primary is back?

There is no leadership election.



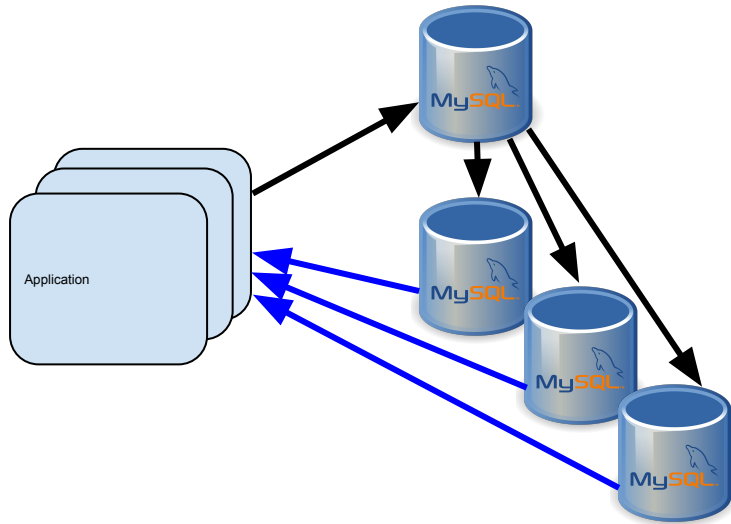
Scalability with MySQL

With MySQL all writes goes first to the primary, then replicated to all others.
All nodes store a complete copy of the data and applies all the changes.



Scalability with MySQL

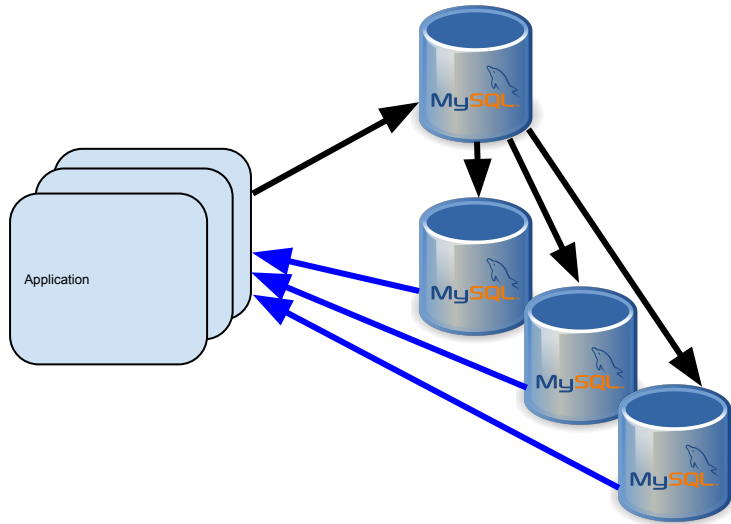
Scaling reads?
Add more replicas



Scalability with MySQL

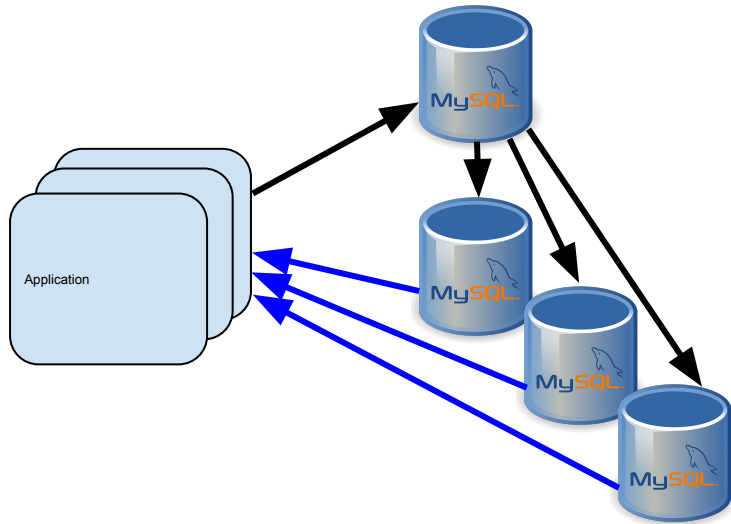
Scaling writes?

Replace the primary with a bigger machine



Scalability with MySQL

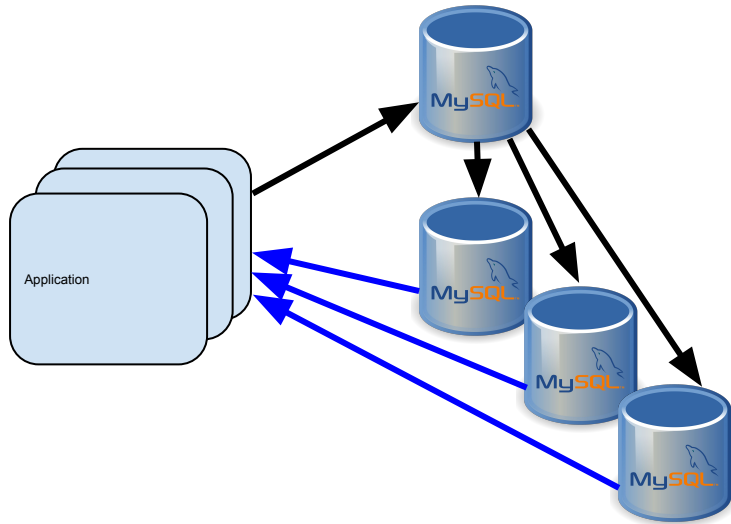
Scaling data volume?
Add bigger disks



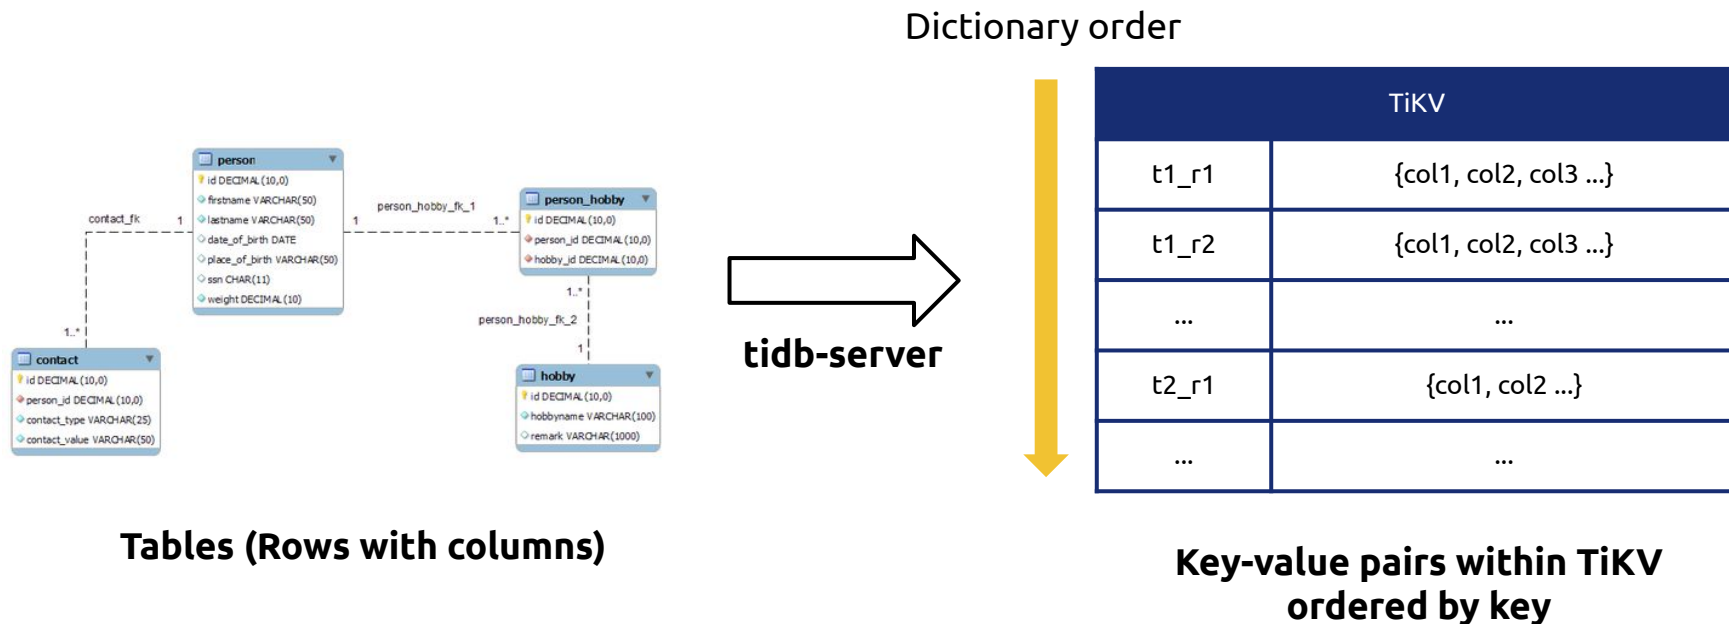
Scalability with MySQL

Need to scale more?

Shard on the application side



Data organization within TiDB



Data organization within TiDB

- For a row in a Table, row data is encoded in key-value pairs like this:

$t\langle\langle\text{tableID}\rangle\rangle_r\langle\langle\text{rowID}\rangle\rangle \Rightarrow [\text{col1}, \text{col2}, \text{col3}, \text{col4}]$

- For unique secondary index:

$t\langle\langle\text{tableID}\rangle\rangle_i\langle\langle\text{indexID}\rangle\rangle_indexedColumnsValue \Rightarrow rowID$

- For non-unique secondary index:

$t\langle\langle\text{tableID}\rangle\rangle_i\langle\langle\text{indexID}\rangle\rangle_indexedColumnsValue_rowID \Rightarrow nil$

Data organization within TiDB

```
CREATE TABLE user_table (  
  id INT,  
  name VARCHAR(64),  
  email VARCHAR(1024),  
  PRIMARY KEY(id)  
);
```

user_table		
1	John	john@example.com
2	Jane	jane@example.com
...

Within TiKV:

t101_r1 => [John,john@example.com]

t101_r2 => [Jane,jane@example.com]

t101_r... => ...

Data organization within TiDB

```
CREATE TABLE user_table (  
  id INT,  
  name VARCHAR(64),  
  email VARCHAR(1024),  
  KEY (name),  
  PRIMARY KEY(id)  
);
```

user_table		
1	John	john@example.com
2	Jane	jane@example.com
...

Within TiKV:

t101_r1 => [John,john@example.com]

t101_r2 => [Jane,jane@example.com]

t101_r... => ...

...

t101_i1_Jane_2 => nil

t101_i1_John_1 => nil

t101_i1_... => nil

Table Partitioning

- Depends on use case, normally TiDB will already improve performance and scalability by distributing the data (can be seen as transparent range sharding/partitioning)
- `CREATE TABLE t (...) PARTITION BY RANGE (col) (partition p1 values less than (1000000),...)`
- Good use cases are for fast delete of large amount of data, instead of a big transactional `DELETE FROM t WHERE col < 1000000`
do a DDL like
`ALTER TABLE t DROP PARTITION p1`
which is as fast and efficient as a drop table, but only for a part[ition] to the table.
- Another use case is to speed up partial scans in TiFlash, like reports for a single range, then range partitioning makes the scanning to only read matching partitions.